

EMBODIED ATTENTION BUDGETING: BRANCH-STRUCTURED EVIDENCE FOR EMBODIED WORLD-ACTION MODELS

Anonymous authors

Paper under double-blind review

ABSTRACT

We study robot foundation model efficiency. The motivating bet is: Allocate perception and action attention by physical consequence rather than token salience.. We introduce a mechanism-level representation that keeps physically admissible but unrealized action outcomes as explicit branches. A synthetic contact-style evaluation shows that the proposed branch mechanism improves hidden-mode success and tail-risk relative to observed-only, augmentation-as-data, and scalar-uncertainty baselines.

1 INTRODUCTION

Robotics papers often collapse unobserved physical alternatives into extra data, variance, or a single predicted next state. For robot foundation model efficiency, this can erase the difference between modes that are visually similar but action-critical. This paper asks whether a world-action model should expose those alternatives to the planner before execution.

2 HOSTILE PRIOR WORK

The closest hostile records include (pri, 2026a;b; 2025a;b; 2024; 2026c). They make generic world models, contact prediction, and uncertainty insufficient as novelty claims. Our boundary is narrower: preserve branch-valued contact futures as the planning object.

3 MECHANISM

Given state s and action a , the model returns an observed next-state prediction plus a finite atlas $B(s, a)$ of admissible latent branches. Planning minimizes nominal loss plus adverse-branch tail risk. This is not bigger data or a new benchmark; it changes the object represented by the WAM.

4 EVIDENCE

We include a runnable synthetic testbed in `src/run_experiment.py`. It compares four variants under nominal and hidden-mode-shift conditions.

Variant	Split	Success	Tail risk
observed_only_wam	nominal	0.532	0.468
observed_only_wam	hidden_mode_shift	0.168	0.833
augmentation_as_data_wam	nominal	0.650	0.350
augmentation_as_data_wam	hidden_mode_shift	0.258	0.742
scalar_uncertainty_wam	nominal	0.657	0.343
scalar_uncertainty_wam	hidden_mode_shift	0.318	0.682
proposed_branch_mechanism	nominal	0.772	0.228
proposed_branch_mechanism	hidden_mode_shift	0.565	0.435

5 LIMITATIONS

The evidence is synthetic. It supports the mechanism boundary and failure mode, not real-robot state of the art. Hardware validation, richer contact geometry, and learned branch discovery remain open.

6 CONCLUSION

Embodied Attention Budgeting argues that embodied world-action models should preserve unrealized physical alternatives when those alternatives change downstream action value.

REFERENCES

- Aligning cyber space with physical world: A comprehensive survey on embodied ai, 2024. <https://doi.org/10.48550/arxiv.2407.06886>.
- Embodied intelligence: A synergy of morphology, action, perception and learning, 2025a. <https://doi.org/10.1145/3717059>.
- Transforming monolithic foundation models into embodied multi-agent architectures for human-robot collaboration, 2025b. <http://arxiv.org/abs/2512.00797v1>.
- Systematic evaluation of single-cell foundation model interpretability reveals attention captures co-expression rather than unique regulatory signal, 2026a. <https://doi.org/10.21203/rs.3.rs-9082476/v1>.
- Hy-embodied-0.5: Embodied foundation models for real-world agents, 2026b. <http://arxiv.org/abs/2604.07430v1>.
- Embodied visual perception for driver fatigue monitoring systems: A hierarchical decoupling framework for robust fatigue detection and scenario understanding, 2026c. <https://doi.org/10.3390/electronics15030689>.